

Lecture 2: Interprocess Communication--Pipes

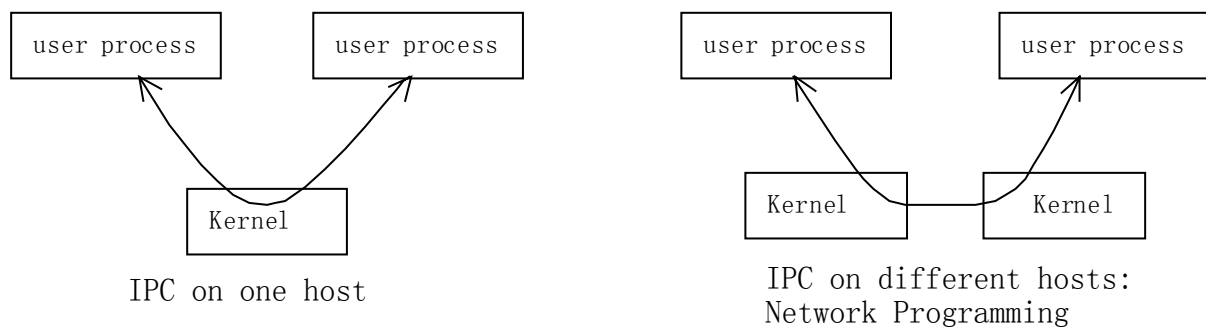
References for Lecture 2:

- 1) Unix Network Programming, W.R. Stevens, 1990, Prentice-Hall, Chapter 3.
- 2) Unix Network Programming, W.R. Stevens, 1999, Prentice-Hall, Chapter 2-6.

Purposes:

Communication is everywhere from intraprocess to Interprocess.

Interprocess communication has 2 forms:



IPC is used for 2 functions:

- 1) Synchronization---Used to coordinate access to resources among processes and also to coordinate the execution of these processes. They are
Record locking,
Semaphores,
Mutexes and Condition variables.
- 2) Message Passing---Used when processes wish to exchange information. Message passing takes several forms such as :
Pipes,
FIFOs,
Message Queues,
Shared Memory.

File or Record Locking:

Used to ensure that a process has exclusive access to a file before using it. Examples: 1) lpr generates a unique sequence number for each print job; see Stevens90:89-91. 2) Couple(husband, wife) share one bank account. [In Unix, record locking is better termed as range locking.](#)

For System V:

```
#include <unistd.h>
```

```
int lockf(int fd, int function, long size);
```

fd---file descriptor(not a file pointer)

size--- define the record size or lock area: [offset, offset + size]. size=0 means the rest of the file. Use lseek() to move the current offset. When the offset position is set to the beginning and size=0 then lock the whole file.

function:

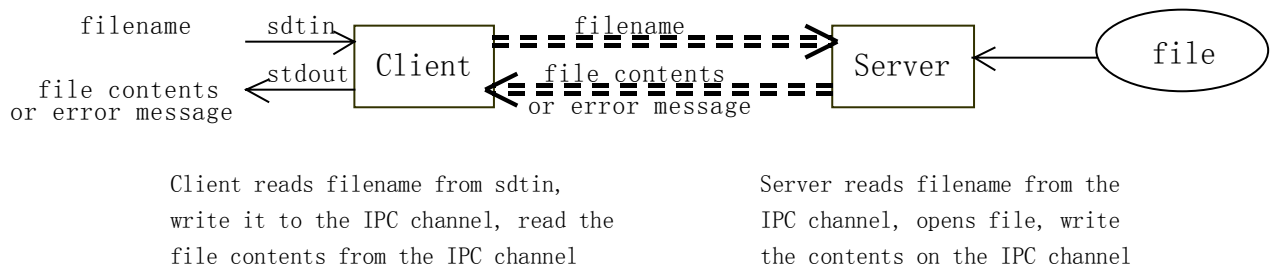
F_ULOCK---unlock a previous lock
 F_LOCK ---lock a region(blocking)
 F_TLOCK ---Test and lock a region(nonblocking)
 F_TEST ---Test a region to see if it is locked.

Example: Use F_TLOCK instead of F_TEST and F_LOCK.

```
...
If (lockf(fd, F_TEST, size) == 0) /* If the region is locked, -1 is returned and the process is in sleep state*/
Re= lockf(fd, F_LOCK, size); /*a small chance that another process locks between TEST and LOCK*/
...
rc=lockf(fd, F_TLOCK, size) /* Test + lock done as an atomic operation, If unsuccessful, lockf()
returns -1 and the calling process continues to do other things*/
```

Atomic operation: one or more operations that are treated as a single operation. No other operation can be executed between the start and end of an atomic operation.

Simple Client-Server or IPC model:



Pipes:

A pipe provides a one-way flow of data.

```
Int pipe (int *filedes);
```

```
Int pipefd[2]; /* pipefd[0] is opened for reading; pipefd[1] is opened for writing */
```

Example to show how to create and use a pipe:

```
main()
{
    int pipefd[2], n;
    char buff[100];
    if (pipe(pipefd) < 0 ) err_sys("pipe error");

    printf("read fd = %d, write fd = %d\n", pipefd[0], pipefd[1]);
    if (write(pipefd[1], "hello world\n", 12) != 12) err_sys("write error");

    if ( (n=read(pipefd[0], buff, sizeof(buff))) <=0) err_sys("read error");
    write(1, buff, n); /*fd=1=stdout*/
}
```

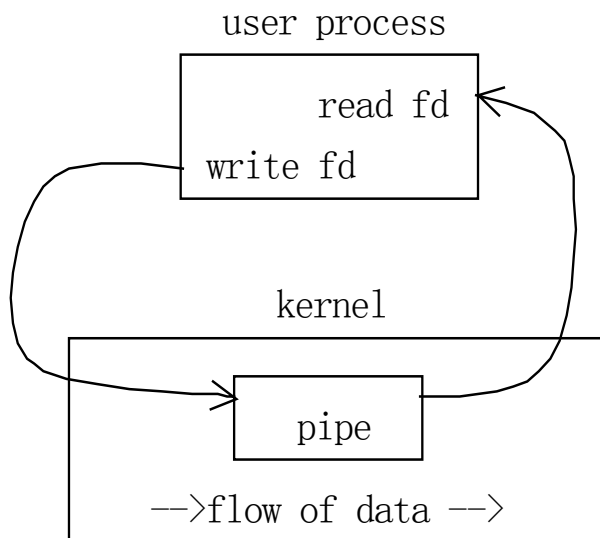
Difference:

- ① { Argument
Parameter
- ② { Call by value
Call by reference
- ③ { File descriptor
File pointer

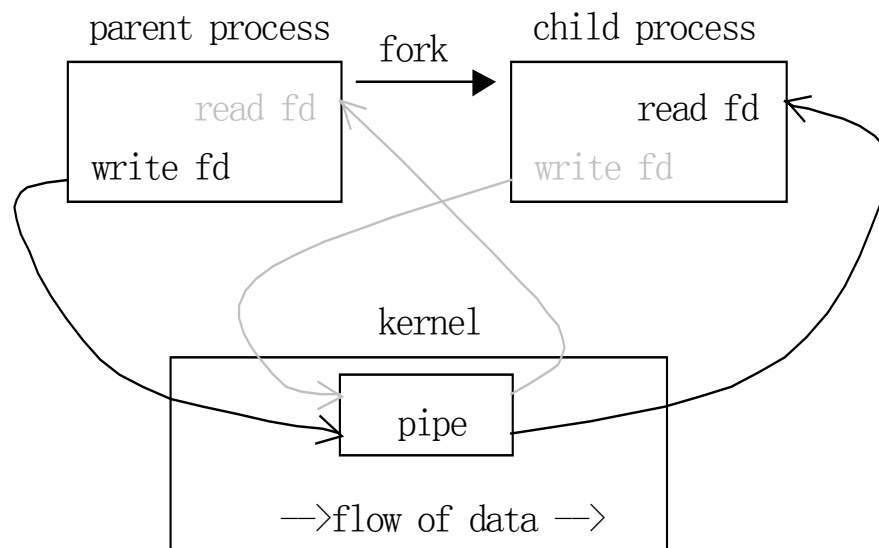
result:

hello world

read fd=3, write df =4

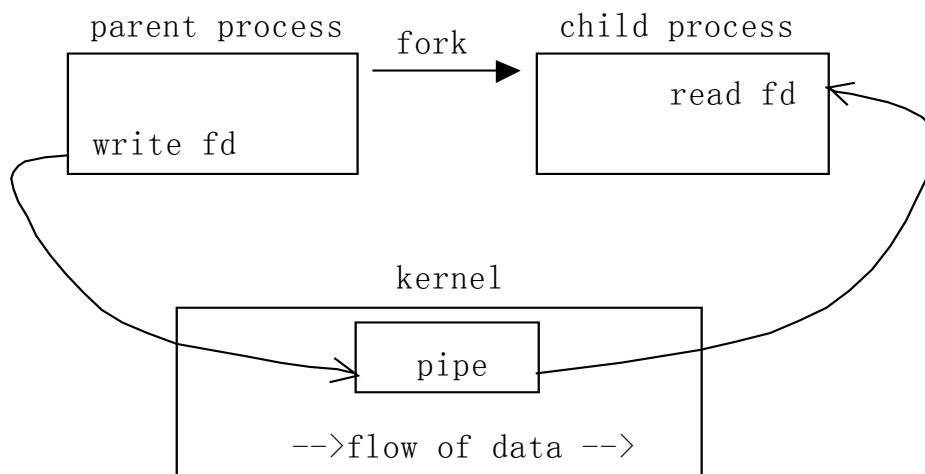


Pipe in a single process



Pipe in a single process, immediately after fork

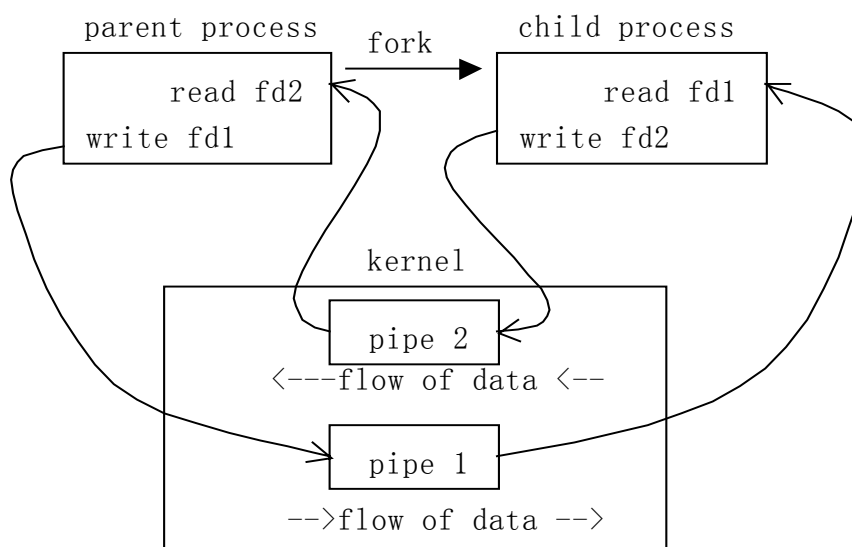
Pipes between two processes: unidirectional



1 pipe between two process: one-way

- Steps :
- 1) opening a pipe
 - 2) forking off another process
 - 3) closing the appropriate pipes on each end

Pipes between two processes: bidirectional



2 pipes to provide a bidirectional flow of data

- Steps:
- 1) create pipe1 + pipe2 : `int pipe1[2], pipe2[2]` -----must be the first step
 - 2) forking off a child process, executing another program as a server
 - 3) parent closes read end of pipe 1 + write end of pipe 2,
 - 4) child closes write end of pipe 1 + read end of pipe 2

Question: Can we use only one pipe to finish bi-directional communications?

Program Example of Simple Client-Server Model:

```
main()
{
    int    childpid, pipe1[2], pipe2[2];

    /* step 1: create pipe1 and pipe2 */
    if (pipe(pipe1) < 0 || pipe(pipe2) < 0)
        err_sys("can't create pipes");

    /*step 2: fork a child process */
    if ( (childpid = fork()) < 0) {
        err_sys("can't fork");
    }

    /* step 3: parent closes read end of pipe 1 and write end of pipe 2*/
    } else if (childpid > 0) {
        close(pipe1[0]);
        close(pipe2[1]);

        client(pipe2[0], pipe1[1]); /*client runs in the parent process*/

        while (wait((int *) 0) != childpid)    /* wait for child */
            ;

        close(pipe1[1]);
        close(pipe2[0]);
        exit(0);
    } else {

        /* step 4: child closes write end of pipe 1 and read end of pipe 2*/
        close(pipe1[1]);
        close(pipe2[0]);

        server(pipe1[0], pipe2[1]); /*server runs in the child process */

        close(pipe1[0]);
        close(pipe2[1]);
        exit(0);
    }
}
```

You may need to replace `err_sys(...)` by `printf(...)` or `perror(...)` for the program to run on your computer.

```

#include    <stdio.h>

#define     MAXBUFF      1024

client(readfd, writefd)
int         readfd;
int         writefd;
{
    char    buff[MAXBUFF];
    int     n;

    /* read the filename from standard input */
    if (fgets(buff, MAXBUFF, stdin) == NULL)
        err_sys("client: filename read error");

    n = strlen(buff);
    if (buff[n-1] == '\n')
        n--;          /* ignore newline from fgets() */

    /* write it to the IPC descriptor, pipe1[1] */
    if (write(writefd, buff, n) != n)
        err_sys("client: filename write error");

    /* read the data from the IPC descriptor, pipe2[0],
       and write to standard output. */
    while ( (n = read(readfd, buff, MAXBUFF)) > 0)
        if (write(1, buff, n) != n) /* fd 1 = stdout */
            err_sys("client: data write error");
    if (n < 0)
        err_sys("client: data read error");
}

```

```

#include    <stdio.h>
#define    MAXBUFF    1024

server(readfd, writefd)
int        readfd;
int        writefd;
{
    char    buff[MAXBUFF];
    char    errmesg[256], *sys_err_str();
    int     n, fd;
    extern int errno;

    /* read the filename from the IPC descriptor, pipe1[0]*/
    if ( (n = read(readfd, buff, MAXBUFF)) <= 0)
        err_sys("server: filename read error");
    buff[n] = '\0';    /* null terminate filename */

    /* open the file from the IPC descriptor, pipe1[0]*/
    if ( (fd = open(buff, 0)) < 0) {
        /* Error.  Format an error message and send it back to the client.    */
        sprintf(errmesg, ": can't open, %s\n", sys_err_str());
        strcat(buff, errmesg);
        n = strlen(buff);
        if (write(writefd, buff, n) != n)
            err_sys("server: errmesg write error");
    } else {

        /* Read the data from the file and write to the IPC descriptor. */
        while ( (n = read(fd, buff, MAXBUFF)) > 0)
            if (write(writefd, buff, n) != n)
                err_sys("server: data write error");
        if (n < 0)
            err_sys("server: read error");
    }
}


```

Properties of Pipe:


- 1) Pipes do not have a name. For this reason, the processes must share a parent process. This is the main drawback to pipes. However, pipes are treated as file descriptors, so the pipes remain open even after fork and exec.
- 2) Pipes do not distinguish between messages; they just read a fixed number of bytes. Newline (\n) can be used to separate messages. A structure with a length field can be used for message containing binary data.
- 3) Pipes can also be used to get the output of a command or to provide input to a command

```
FILE *popen(const char *command, const char *type);  
Int pclose(FILE *stream);
```

When type is “r”:



When type is “w”:



For example:

```
#include <stdio.h>  
#define MAXLINE 1024  
main()  
{ int n;  
  char line[MAXLINE];  
  FILE *fp;  
  
  fp=popen("cat .cshrc", "r");  
  
  \*read the lines in .cshrc from fp*\br/>  while ((fgets(line, MAXLINE, fp)) != NULL) {  
    n=strlen(line);  
    if (write(1, line, n)!=n) printf("print data error");  
    pclose(fp);  
  }  
}
```

Notes:

- Please notice the difference between fgets(...) and read(...). See our website FAQ.
- Not every Unix command has output such as mv, cp and In such cases, fgets() and read() will return NULL when reading from fp=popen(...).
- **cd** is a special Unix command. You cannot use popen("cd","r") or system("cd").